

Computing Infrastructures - Notes

260236

August 2024

Preface

Every theory section in these notes has been taken from two sources:

- The Datacenter as a Computer: Designing Warehouse-Scale Machines, Third Edition. [1]
- Quantitative System Performance: Computer System Analysis Using Queueing Network Models. [2]

About:

 [GitHub repository](#)

Contents

1	Introduction: definition of Data Center and Computing Infrastructure	4
2	Hardware Infrastructures	5
2.1	System-level	5
2.1.1	Computing Infrastructures and Data Center Architectures	5
2.1.1.1	Overview of Computing Infrastructures	5
2.1.1.2	The Datacenter as a Computer	10
2.1.1.3	Warehouse-Scale Computers	11
2.1.1.4	Multiple Data Centers	12
2.1.1.5	Warehouse-Scale Computing / Data Centers Availability	13
2.1.1.6	Architectural overview of Warehouse-Scale Computing	13
2.2	Node-level	15
2.2.1	Server (computation, HW accelerators)	15
2.2.1.1	Rack Servers	17
2.2.1.2	Blade Servers	18
2.2.1.3	Machine Learning	19
2.2.2	Storage (type, technology)	22
2.2.2.1	Files	23
2.2.2.2	HDD	26
2.2.2.3	SSD	29
2.2.2.4	RAID	37
2.2.3	Networking (architecture and technology)	51
3	Software Infrastructure	52
3.1	Virtualization	52
3.2	Computing Architectures	53
4	Methods	54
4.1	Reliability and availability of datacenters	54
4.2	Disk performance	55
4.2.1	HDD	55
4.3	Scalability and performance of datacenters	60
	Index	62

1 Introduction: definition of Data Center and Computing Infrastructure

There's no single definition of a Data Center, but it can be summarized as follows.

Definition 1

Data Centers are buildings where multiple servers and communication gear are co-located because of their common environmental requirements and physical security needs, and for ease of maintenance. [1]

Definition 2

A **Computing Infrastructure** (or IT Infrastructure) is a technological infrastructure that provides hardware and software for computation to other systems and services.

Traditional data centres have the following characteristics:

- **Host a large number** of relatively small or medium sized **applications**;
- Each **application is running on a dedicated HW infrastructure** that is de-coupled and protected from other systems in the same facility;
- **Applications tend not to communicate each other.**

Those **data centers host hardware and software for multiple organizational units** or even **different companies**.

2 Hardware Infrastructures

2.1 System-level

2.1.1 Computing Infrastructures and Data Center Architectures

2.1.1.1 Overview of Computing Infrastructures

A number of computing infrastructures exist:

- **Cloud** offers virtualized computing, storage and network resources with highly-elastic capacity.
- **Edge Servers** are on-premises hardware resources that perform more compute-intensive data processing.

In other words, an edge server is a piece of hardware that performs data computation at the end (or “edge”) of a network. Like a regular server, an edge server can provide compute, networking, and storage functions.¹

- **IoT and AI-enabled Edge Sensors** are hardware devices where the data acquisition and partial processing can be performed at the edge of the network.



Figure 1: An **example** of Computing Infrastructures. [4]

The **Computing Continuum**, a novel paradigm that extends beyond the current silos of cloud and edge computing, can enable the seamless and dynamic deployment of applications across diverse infrastructures. [3]

¹More info [here](#).



Figure 2: The Computing Continuum. [4]

In the following pages, we analyze the computing infrastructures mentioned in the previous example.

Data Centers

The definition of a Data Centers can be found on page 4.

✓ Data Centers Advantages

- Lower IT costs.
- High Performance.
- Instant software updates.
- “Unlimited” storage capacity.
- Increased data reliability.
- Universal data access.
- Device Independence.

🚫 Data Centers Disadvantages

- Require a constant internet connection.
- Do not work well with low-speed connections.
- Hardware Features might be limited.
- Privacy and security issues.
- High power Consumption.
- Latency in taking decision.

Internet-of-Things (IoT)

An **Internet of Things (IoT)** device is any everyday object embedded with sensors, software, and internet connectivity.

This allows to collect and exchange data with other devices and systems, typically over the internet, with limited need of process and store data.

Some **examples** are [Arduino](#), [STM32](#), [ESP32](#), [Particle Argon](#).

✓ Internet-of-Things Advantages

- Highly Pervasive.
- Wireless connection.
- Battery Powered.
- Low costs.
- Sensing and actuating.

🚫 Internet-of-Things Disadvantages

- Low computing ability.
 - Constraints on energy.
 - Constraints on memory (RAM/FLASH).
 - Difficulties in programming.
-

Embedded (System) PCs

An **Embedded System** is a computer system, a combination of a computer processor, computer memory, and input/output peripheral devices, that has a dedicated function within a larger mechanical or electronic system.

A few **examples**: [Odroid](#), [Raspberry](#), [jetson nano](#), [Google Coral](#).

✓ Embedded System Advantages

- Persuasive computing.
- High performance unit.
- Availability of development boards.
- Programmed as PC.
- Large community.

🚫 Embedded System Disadvantages

- Pretty high power consumption.
- (Some) Hardware design has to be done.

Edge/Fog Computing Systems

The key **difference** between **Fog Computing** and **Edge Computing** is associated with the location **where the data is processed**:

- In **edge computing**, the data is processed closest to the sensors.
- In **fog computing**, the computing is moved to processors linked to a local area network (IoT gateway).

Edge computing places the intelligence in the connected devices themselves, whereas, fog computing puts in the local area network.



✓ Fog/Edge Advantages

- High computational capacity.
- Distributed computing.
- Privacy and security.
- Reduced Latency in making a decision.

🔴 Fog/Edge Disadvantages

- Require a power connection.
- Require connection with the Cloud.

Feature	Edge Computing	Fog Computing
Location	Directly on device or nearby device.	Intermediary devices between edge and cloud.
Processing Power	Limited due to device constraints, sending data to central server for analysis.	More powerful than edge devices. However, sending data to a central server for analysis.
Primary Function	Real-time decision-making, low latency. However, central server analyzing combined data and sending only relevant information further.	Pre-process and aggregate data, reduce bandwidth usage. However, central server analyzing combined data and sending only relevant information further.
Advantages	Low latency, reduced reliance on cloud, security for sensitive data.	Bandwidth efficiency, lower cloud costs, complex analysis capabilities.
Disadvantages	Limited processing power, single device focus.	Increased complexity, additional infrastructure cost.

Table 1: Differences between Edge and Fog Computing Systems.

2.1.1.2 The Datacenter as a Computer

In the last few decades, computing and storage have moved from PC-like clients to smaller, often mobile, devices combined with extensive internet services. Furthermore, traditional enterprises are also shifting to Cloud computing.

The **advantages** of this migration are:

- **User-side:**
 - **Ease of management** (no configuration or backups needed);
 - The **availability of the service is everywhere**, but we need connectivity.
- **Vendors-side:**
 - **SaaS (Software-as-a-Service)** allows **faster application development** (more accessible to make changes and improvements);
 - **Improvements and fixes** in the software are **more straightforward inside their data centers** (instead of updating many millions of clients with peculiar hardware and software configurations);
 - The **hardware deployment** is restricted to a **few well-tested configurations**.
- **Server-side:**
 - **Faster introduction of new hardware devices** (e.g., HW accelerators or new hardware platforms);
 - Many application **services can run at a low cost per user**.

Finally, another advantage is that **some workloads require so much computing capability that they are a more natural fit in the datacenter** (and not in client-side computing). For example, the search services (web, images, and so on) or the Machine and Deep Learning.

2.1.1.3 Warehouse-Scale Computers

The trends toward server-side computing and widespread internet services created a new class of computing systems: **Warehouse-Scale Computers**.

Definition 1

Warehouse-Scale Computers (WSCs) is intended to draw attention to the most distinctive feature of these machines: **the massive scale of their software infrastructure, data repositories and hardware platform**.

? What is a *program* at a WSC?

In Warehouse-Scale Computing **the program is an internet service**, which may **consist of tens or more individual programs that interact to implement complex end-user services** such as *email*, *search*, or *maps*. These programs might be implemented and maintained by different teams of engineers, perhaps even across organizational, geographic, and company boundaries.

⚖ Difference between WSCs and Data Centers

WSCs currently power the services offered by companies such as Google, Amazon, Microsoft, and others. The main difference from traditional data centers (see more on page 4) is that **WSCs belong to a single organization, use a relatively homogeneous hardware and system software platform, and share a common systems management layer**. In contrast with the typical data center that belongs to multiple organizational units or even different companies, use dedicated HW infrastructure in order to run a large number of applications (more details on page 4).

? How is the WSC organized?

The **software on WSCs**, such as Gmail, runs on a scale far beyond a single machine or rack: **it runs on clusters of hundreds to thousands of individual servers**. Therefore, the machine, the computer, is itself this **large cluster or aggregation of servers** and must be **considered a single computing unit**.

Most importantly, WSCs run fewer vast applications (internet services). An **advantage** is that the **shared resource management infrastructure allows significant deployment flexibility**. Finally, the requirements of:

- **Homogeneity**
- **Single-Organization Control**
- **Cost Efficiency**

Motivate designers to take new approaches to constructing and operating these systems.

2.1.1.4 Multiple Data Centers

Sometimes the data centers are **located far apart**. **Multiple data centers** are (often) **replicas of the same service**:

- To *reduce user latency*
- To *improve service throughput*

Typically, a request is fully processed within one data center.

The world is divided into **Geographic Areas (GAs)**. Each Area is defined by Geo-political boundaries (or country borders). Also, there are at least two computing regions in each geographical Area.

The **Computing Regions (CRs)** are the smallest geographic unit of the infrastructure from the customer's perspective. Multiple Data centers within the same region are not exposed to customers.

However, they are defined by a latency-defined perimeter, typically less than 2ms for round-trip latency. Finally, they're located hundreds of miles apart, with considerations for different flood zones, etc. It is too far from synchronous replication but suitable for disaster recovery.

The **Availability Zones (AZs)** are finer-grain **locations within a single computing region**. They allow customers to run mission-critical applications with high availability and fault tolerance to Data Center failures. Because there are fault-isolated locations with redundant power, cooling, and networking (they are different from the concept of the Availability Set).

This hierarchical structure ensures efficient data management and compliance with local data laws while optimizing network performance through strategically placing data centers.

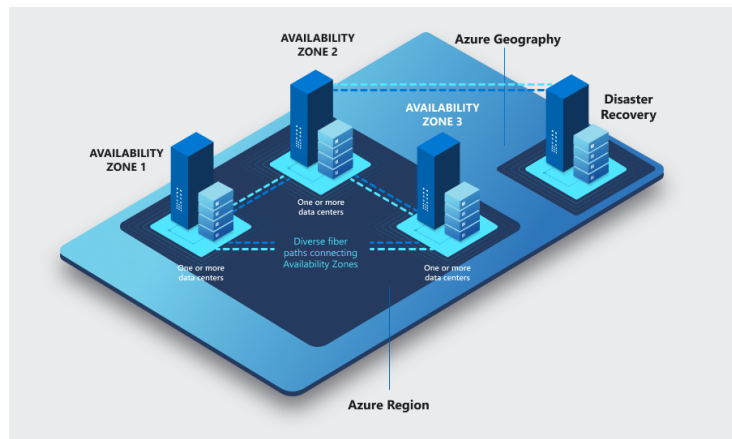


Figure 3: Example of [Azure Availability Zones](#).

2.1.1.5 Warehouse-Scale Computing / Data Centers Availability

The services provided through WSCs (or DCs) **must guarantee high availability**, typically aiming for at least 99.99% uptime (e.g. one hour of downtime per year).

Some **examples**:

- 99,90% on single instance VMs with premium storage for a more accessible lift and shift;
- 99,95% VM uptime SLA for Availability Sets (AS) to protect for failures within a data center;
- 99,99% VM uptime SLA through Availability Zones.

Such fault-free operation is more accessible when an extensive collection of hardware and system software is involved.

WSC workloads must be designed to gracefully tolerate large numbers of component faults with little or no impact on service level performance and availability!

2.1.1.6 Architectural overview of Warehouse-Scale Computing

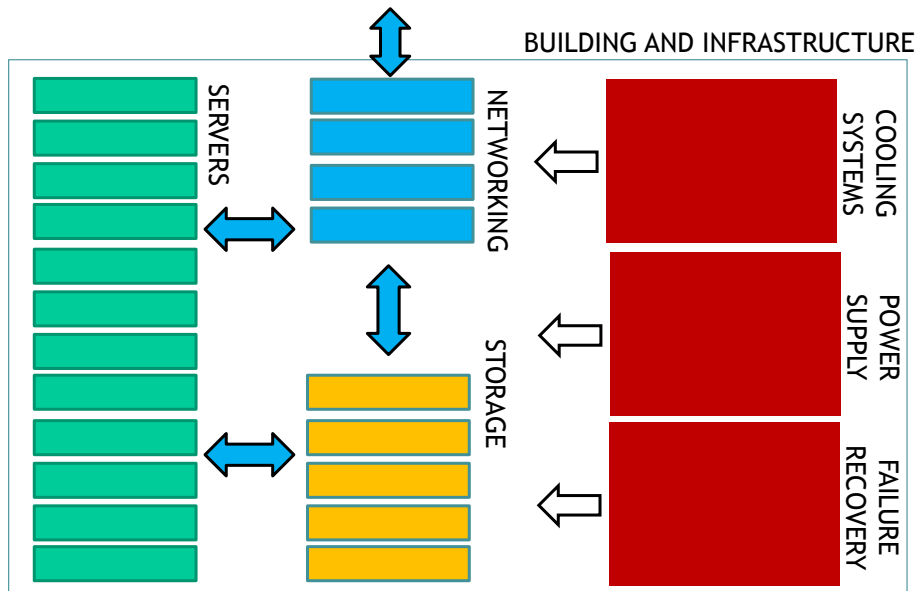


Figure 4: Architectural overview of Warehouse-Scale Computing.

- **Server** (section 2.2.1, page 15). Servers are the **leading processing equipment**: different types according to CPUs, RAM, local storage, accelerators, and form factor. The servers are **hosted on individual shelves** and are the **basic building blocks of Data Centers and Warehouse-Scale Computers**. They are interconnected by hierarchies of networks and supported by the shared power and cooling infrastructure.
- **Storage** (section 2.2.2, page 22). Disks, flash SSDs, and Tapes are the **building blocks** of today's **WSC storage systems**. These devices are **connected to the Data Center network and managed by sophisticated distributed systems**.

Some **examples**:

- Direct Attached Storage (DAS)
 - Network Attached Storage (NAS)
 - Storage Area Networks (SAN)
 - RAID controllers
- **Networking** (section 2.2.3, page 51). The **Data Center Network (DCN)** **enables efficient data transfer and interaction between various components**. The data processing ecosystem within the DCs needs to reach the DC services from outside. Communication equipment includes switches, Routers, cables, DNS or DHCP servers, Load balancers, Firewalls, etc.
 - **Building and Infrastructure**. WSC has other essential components related to power delivery, cooling, and building infrastructure that must be considered. Some interesting numbers:
 - Data Centers with up to 110 football-pitch size.
 - 2-100s MW power consumption (100k houses), and the largest in the world is 650 MW.

2.2 Node-level

2.2.1 Server (computation, HW accelerators)

Servers are like ordinary PCs, usually more powerful, but with a **form factor that allows them to fit into the shelves** (such as rack, blade enclosure format, or tower; the differences are explained later). They are usually built in a tray or blade enclosure format, **housing the motherboard, the chipset, and additional plug-in components**.

The **motherboard** acts as the central hub, **connecting all the crucial components of the server and enabling them to communicate and work together**.

It provides sockets and plug-in slots to install CPUs, memory modules (DIMMs), local storage (such as Flash SSDs or HDDs), and network interface cards (NICs) to satisfy the range of resource requirements.

The **chipsets** and **additional components** are grouped in the following way:

- Number and type of CPUs:
 - From 1 to 8 CPU socket.
 - Intel Xeon Family, AMD EPYC, etc.
- Available RAM:
 - From 2 to 192 DIMM Slots.
- Locally attached disks:
 - From 1 to 24 Drive Bays.
 - HDD or SSD.
 - SAS (higher performance but more expensive) or SATA (for entry-level servers).
- Other special purpose devices:
 - From 1 to 20 GPUs per node, or TPUs.
 - NVIDIA Pascal, Volta, etc.
- Form factor:
 - From 1 unit to 10 units.
 - Tower.

Differences between Rack, Blade and Tower

Tower Server

A **Tower Server** looks and feels much like a **traditional** tower **PC**.

✓ Advantages

- ✓ **Scalability and ease of upgrade.** Customized and upgraded based on necessity.
- ✓ **Cost-effective.** Tower servers are probably the *cheapest of all kinds of servers*.
- ✓ **Cools easily.** Since a tower server has a low overall component density, it cools down easily.

🚫 Disadvantages

- ✗ **Consumes a lot of space.** These servers are difficult to manage physically.
- ✗ **Provides a basic level of performance.** A tower server is *ideal for small business that have a limited number of clients*.
- ✗ **Complicated cable management.** Devices aren't easily routed together.

2.2.1.1 Rack Servers

Rack Servers are unique **shelves that accommodate all the IT equipment** and allow their interconnection. The racks are used to store these rack servers.

Server racks are measured in **Rack Units**, or "U". One U is approximately 44.45 millimeters. The main advantage of these racks is that they **allow designers to stack up other electronic devices and servers**.

A rack server is designed to be positioned in a bay by vertically stacking servers one over the other along with other devices (storage units, cooling systems, network peripherals, and batteries).

✓ Advantages

- ✓ **Failure containment.** Very little effort to identify, remove, and replace a malfunctioning server with another.
- ✓ **Simplified cable management.** Easy and efficient to organize cables.
- ✓ **Cost-effective.** Computing power and efficiency at relatively lower costs.

✗ Disadvantages

- ✗ **Power usage.** Needs of additional cooling systems due to their high overall component density, thus consuming more power.
- ✗ **Maintenance.** Since multiple devices are placed in racks together, maintaining them gets considerably tough with the increasing number of racks.

2.2.1.2 Blade Servers

Blade Servers are the latest and the most advanced type of servers in the market. They can be termed hybrid rack servers, where servers are placed inside blade enclosures, forming a blade system.

The **most significant advantage** of blade servers is that these servers are the **most minor types of servers available now** and are **great for conserving space**.

Finally, a blade system also meets the IEEE standard for rack units, and each rack is measured in the units of "U".

✓ Advantages

- ✓ **Size and form factor.** They are smallest and the most compact servers, requiring minimal physical space. Blade servers offer *higher space efficiency* compared to traditional rack-mounted servers.
- ✓ **Cabling.** Blade server don't involve the cumbersome tasks of setting up cabling. Although you still might have to deal with the cabling, it is near to negligible when compared to tower and rack servers.
- ✓ **Centralized management.** Blade enclosures typically come with centralized management tools that allow administrators to easily monitor, configure and update all blades from a single interface.
- ✓ **Load balancing, failover, scalability.** Uniform system, shared components (including network), simple addition/removal of servers.

✗ Disadvantages

- ✗ **Expensive configuration and Higher initial cost.** Although upgrading the blade server is easy to handle and manage, the initial configuration or the setup requires more effort and higher initial investment.
- ✗ **Vendor Lock-In.** Blade servers typically require the use of the manufacturer's specific blades and enclosures, leading to vendor lock-in. This can limit flexibility and potentially increase costs in the long run.
- ✗ **Cooling.** Blade servers come with high component density. Therefore, special accommodations have to be arranged for these servers to ensure they don't get overheated. Heating, ventilation, and air conditioning systems (HVAC) must be carefully managed and designed.

2.2.1.3 Machine Learning

Deepening: Machine Learning (supervised learning)

Machine learning (ML) is a branch of artificial intelligence (AI) and computer science that focuses on the using data and algorithms to enable AI to imitate the way that humans learn, gradually improving its accuracy ([source](#)).

[UC Berkeley](#) breaks out the learning system of a machine learning algorithm into three main parts:

1. A Decision Process: In general, machine learning algorithms are used to make a prediction or classification. Based on some input data, which can be labeled or unlabeled, your algorithm will produce an estimate about a pattern in the data.
2. An Error Function: An error function evaluates the prediction of the model. If there are known examples, an error function can make a comparison to assess the accuracy of the model.
3. A Model Optimization Process: If the model can fit better to the data points in the training set, then weights are adjusted to reduce the discrepancy between the known example and the model estimate. The algorithm will repeat this iterative “evaluate and optimize” process, updating weights autonomously until a threshold of accuracy has been met.

The main goal is to learn a target function that can be used for prediction. Given a training set of labeled examples $\{(x_1, y_1), \dots, (x_n, y_n)\}$, estimate the prediction function f by minimizing the prediction error on the training set:

$$y = f(x)$$

Where y is the output, f is the prediction function and the x is an image feature.

Deepening: Artificial Neural Network

The **Artificial Neural Network** is a computational model inspired by the human brain (perceptron). It consists of interconnected nodes (neurons) organized in layers to process and analyze data and used to learn data representation from data (learn features and the classifier/regressor).

The learning process of a Neural Network is as follows: Neurons make decisions (activation functions). There are wights, so the connections between neurons are strengthened or weakened through training- randomly initialized.

The (training data) Neural Networks learn from historical data and ex-

amples. Then, labeled data are provided.

Deepening: effects of ML and ANN

Deep learning models began to appear and be widely adopted, enabling specialized hardware to power a broad spectrum of machine learning solutions.

Since 2013, AI learning compute requirements have doubled every 3.5 months (vs. 18-24 months expected from [Moore's Law](#)).

To satisfy the growing compute needs for deep learning, **WSCs deploy specialized accelerator hardware:**

- Graphical Processing Units (GPUs) are used for data-parallel computations (the same program is executed on many data elements in parallel). In order to use parallel programming, high-level languages such as CUDA, OpenCL, OPENACC, OpenMP, and SYCL exist. This technique allows up to 1000x faster than CPU.
- Tensor Processing Unit (TPU), where Tensor is a n-dimensional matrix, are used for training and inference.
- Field-Programmable Gate Array (FPGA) are programmable hardware devices. The device user can program an array of logic gates ("configured") in the field instead of the people who designed it. An array of carefully designed and interconnected digital subcircuits that efficiently implement common functions, offering very high levels of flexibility. The digital subcircuits are called configurable logic blocks (CLBs).

FPGA Applications in Data Centers:

- Network acceleration: FPGAs can offload specific processing tasks from CPUs, improving overall network performance and reducing CPU workload.
- Security acceleration: Encryption, decryption, and other security-related tasks can be accelerated using FPGAs, enhancing data centre security while maintaining performance.
- Data analytics: FPGAs can accelerate specific algorithms in data analytics workloads, leading to faster data processing and analysis.
- Machine learning: FPGAs can be configured to efficiently implement specific machine learning algorithms, potentially offering performance advantages for specialized tasks.

Advantages	Disadvantages
CPU	
<ul style="list-style-type: none"> • Easy to be programmed and support any programming framework. • Fast design space exploration and run your applications. 	<ul style="list-style-type: none"> • Suited only for simple AI models that do not take long to train and for small models with small training set.
GPU	
<ul style="list-style-type: none"> • Ideal for applications in which data need to be processed in parallel like the pixels of images or videos. 	<ul style="list-style-type: none"> • Programmed in languages like CUDA and OpenCL and therefore provide limited flexibility compared to CPUs.
TPU	
<ul style="list-style-type: none"> • Very fast at performing dense vector and matrix computations and are specialized on running very fast programming based on Tensorflow. 	<ul style="list-style-type: none"> • For applications and models based on the TensorFlow. • Lower flexibility compared to CPUs and GPUs.
FPGA	
<ul style="list-style-type: none"> • Higher performance, lower cost and lower power consumption compared to other options like CPUs and GPU. 	<ul style="list-style-type: none"> • Programmed using OpenCL and High-Level Synthesis (HLS). • Limited flexibility compared to other platforms.

2.2.2 Storage (type, technology)

Data has significantly grown in the last few years due to sensors, industry 4.0, AI, etc. The growth favours the **centralized storage strategy** that is focused on the following:

- **Limiting redundant data**
- **Automatizing replication and backup**
- **Reducing management costs**

The *storage technologies* are many. One of the oldest but still used is the **Hard Disk Drive (HDD)**, a magnetic disk with mechanical interactions. Due to its mechanical movement, the **solid-state drive (SSD)** is the best solution (quality-price) because there are no mechanical or moving parts, and they are built out of transistors (NAND flash-based devices). The **non-volatile memory express (NVMe)** also exists, which is the **latest industry standard** for running PCIe² SSDs.

As for price classification, we can see that the NVMe is the most expensive solution:

1. NVMe (between 100€ and 200€ for 1TB)
2. SSD (between 70€ and 100€ for 1TB)
3. HDD (between 40€ and 60€ for 1TB)

For these reasons, it is reasonable to use a **hybrid solution** (HDD + SSD):

- A speed storage technology (**SSD or NVMe**) as **cache** and **several HDDs for storage**. It is a combination used by some servers: a small SSD with a large HDD to have a faster disk.
- Some HDD manufacturers produce Solid State Hybrid Disks (SSHD) that combine a small SSD with a large HDD in a single unit.

²**PCIe (peripheral component interconnect express)**. is an interface standard for connecting high-speed components

2.2.2.1 Files

An OS can see the disks as a collection of **data blocks** that can be read or written independently. To allow the ordering/management among them, **each block is characterized by a unique numerical address called LBA (Logical Block Address)**. Typically, the OS groups blocks into clusters³ to simplify the access to the disk. Typical cluster sizes range from 1 disk sector (512 B, or 4 KB) to 128 sectors (64 KB).

Each *cluster* contains:

- **File data.** The actual content of the files.
- **Metadata.** The information required to support the file system:
 - File names
 - Directory structures and symbolic links
 - Creation, modification, and last access dates
 - Security information (owners, access list, encryption)
 - **Links to the LBA where the file content can be located on the disk**

The disk can thus contain **several types of clusters**:

- Metadata:
 - Fixed position (to bootstrap the entire file system)
 - Variable position (to hold the folder structure)
- File data (the actual content of the files)
- Unused space (available to contain new files and folders)




Figure 5: A cluster can be seen visually as an array. In this image, for example, we've shown three types of cluster: metadata fixed position (blue), metadata variable position (green), file data (orange), unused space (white).


³**Clusters** are the minimal units an OS can read from or write to a disk.

The following explanation introduces some basic operations on the files to see what happens inside the disks.

- **Reading.** To read a file:
 1. Access the metadata, variable position (because it contains the folder structure), to locate its block;
 2. Access the blocks to read the contents of the file.



- **Writing.** To write a file:
 1. Access the metadata, variable position (because it contains the folder structure), to find free space.
 2. Write the data in the allocated blocks (cluster type: unused space).



Since the *file system can only access clusters*, the **actual space taken up by a file on a disk is always a multiple of the cluster size**. Given:

- s , the *file size*
- c , the *cluster size*

Then the **actual size on the disk a** can be calculated as:

$$a = \text{ceil} \left(\frac{s}{c} \right) \times c \quad (1)$$

Where ceil rounds a number up to the nearest integer. It's also possible to calculate the **amount of disk space wasted by organising the file into clusters (wasted disk space w)**:

$$w = a - s \quad (2)$$

A formal way to refer to wasted disk space is **internal fragmentation** of files.

Example 1: internal fragmentation

- File size: 27 byte
- Cluster size: 8 byte

The *actual size* on the disk is:

$$a = \text{ceil}\left(\frac{27}{8}\right) \cdot 8 = \text{ceil}(3.375) \cdot 8 = 4 \cdot 8 = 32 \text{ byte}$$


And the internal fragmentation w is:

$$w = 32 - 27 = 5 \text{ byte}$$

- **Deleting.** To delete a file:

1. Just update the metadata, variable position (because it contains the folder structure), to say that the blocks where the file was stored are no longer used by the OS.


Deleting a file never actually deletes the data on the disk: if a new file is written to the same clusters, the old data is replaced by the new.



- **External fragmentation.** As the disk's life evolves, there might **not be enough space to store a file contiguously**.

In this case, the file is split into smaller chunks and inserted into the free clusters spread over the disk.

The effect of **splitting a file into non-contiguous clusters** is called **external fragmentation**.



2.2.2.2 HDD

A **Hard Disk Drive (HDD)** is a data storage device that uses rotating disks (platters) coated with magnetic material.

Data is read randomly, meaning individual data blocks can be stored or retrieved in any order rather than sequentially.

An HDD consists of one or more rigid (*hard*) rotating disks (platters) with magnetic heads arranged on a moving actuator arm to read and write data to the surfaces.




Figure 6: Hard Drive Disk anatomy.

Externally, hard drives expose a large number of **sectors** (blocks):

- Typically, 512 or 4096 bytes.
- Individual **sector writes** are **atomic**.
- Multiple sectors write it may be interrupted (**torn write**⁴).

The geometry of the drive:

- The sectors are arranged into **tracks**.
- A **cylinder** is a particular track on multiple platters.
- Tracks are arranged in concentric circles on **platters**.
- A disk may have multiple double-sided platters.

The **driver motor spins the platters at a constant rate**, measured in **Revolutions Per Minute (RPM)**.

⁴Torn writes happen when only part of a multi-sector update is written successfully to disk.




Figure 7: Example of HDD geometry.

Given the architecture of the HDD, there are **four types of delay**:

- **Rotational Delay** is the **time to rotate the desired sector to the read head**, and it's related to RPM.
- **Seek Delay** is the **time to move the read head to a different track**.
- **Transfer time** is the **time to read or write bytes**.
- **Controller Overhead** is the **overhead for the request management**.

In order to reduce the delay in the HDD, the companies use a high-speed and tiny memory (8, 16 or 32 MB) called **cache** (also called **track buffer**). The cache is used when there is a:

- **Read caching**. It reduces read delays due to seeking and rotation.
- **Write caching**. It is divided into two different implementations:
 - **Write Back cache**: the drive reports that writes are complete after they have been cached. The disadvantage is that it has an inconsistent state if the power goes off before the write-back event.
 - **Write Through cache**: the drive reports that writes are complete after they have been written to disk.

So, the **caching helps improve disk performance**. The critical idea under caching is that if there is a **queue of requests to the disk**, they can be **reordered to improve performance**. The estimation of the request length is feasible, knowing the position of the data on the disk.

There are several **scheduling algorithms**:

- **First Come, First Serve (FCFC)**. It is the most basic scheduler, serving requests in order. The *disadvantage* is that there is much time spent seeking.
- **Shortest Seek Time First (SSTF)**. Its primary purpose is to minimize seek time by always selecting the block with the shortest seek time. The *advantage* is that it is optimal and can be easily implemented. The main *disadvantage* is that it is prone to starvation.
- **SCAN**, otherwise known as the Elevator Algorithm. The head sweeps across the disk, servicing requests in order. The *advantage* is that it performs reasonably well and does not suffer starvation. The *disadvantage* is that the average access times are higher for requests at high and low addresses.
- **C-SCAN (Circular SCAN)**. It is like the SCAN algorithm, but only service requests in one direction. The *advantage* is fairer than SCAN. However, the *disadvantage* is that it has worse performance than SCAN.
- **C-LOOK**. It is a C-SCAN variant that peeks at the upcoming addresses in the queue. The head only goes as far as the last request.

2.2.2.3 SSD

The **solid-state drive (SSD)** does not have mechanical or moving parts like an HDD. It is built out of transistors (like memory and processors). It has higher performance than HDD.

It stores bits in cells. Each cell can have:

- Single-Level Cell (SLC), a single bit per cell.
- Multi-Level Cell (MLC), two bits per cell.
- Triple-Level Cell (TLC), three bits per cell.
- And so on... QLC, PLC, etcetera.

Internally, the SSD has a lot of NAND flashes, which are organized into Pages and Blocks. Some terminology:

- A **Page** contains **multiple logical block** (e.g. 512 B - 4 KB) **addresses** (LBAs). It is the **smallest unit that can be read/written**. It is a sub-unit of an erase block and consists of the number of bytes which can be read/written in a single operation. The states of each page are:
 - **Empty (ERASED)**: it does not contain data.
 - **Dirty (INVALID)**: it contains data, but this data is no longer in use (or never used).
 - **In use (VALID)**: the page contains data that can be read.
- A Block (or **Erase Block**) typically consists of **multiple pages** (e.g. 64) with a total capacity of around 128-256 KB. It is the **smallest unit that can be erased**.

When passing from the HDD to SSD, there is a problem known as Write Amplification (WA). **Write amplification (WA)** is an **undesirable phenomenon associated with flash memory and solid-state drives (SSDs)** where the actual amount of information physically written to the storage media is a multiple of the logical amount intended to be written.


Example 2

Given a hypothetical SSD:

- Page Size: 4 KB
- Block Size: 5 Pages
- Drive Size: 1 Block
- Read Speed: 2 KB/s
- Write Speed: 1 KB/s

Let us write a 4 KB text file to the brand-new SSD. The overall writing time is 4 seconds (write speed \times file dimension, $1 \text{ KB/s} \times 4 \text{ KB}$).

Now, let us write an 8 KB picture file for the almost brand-new SSD; thankfully, there is space. The overall Writing time is 8 seconds, and the calculation is the same as above.



Now, consider that the first file inserted on the first page is unnecessary.

Finally, let's write a 12 KB pic to the SSD. Theoretically, the image should take 12 seconds. However, it is wrong! The SSD has only two empty pages and one dirty page (invalid). Then, the operations are:

1. Read block into cache.
2. Delete page from cache (set dirty page).
3. Write a new picture into the cache.
4. Erase the old block on the SSD.
5. Write cache to SSD.

The OS only thought it was writing 12 KBs of data when the SSD had to read 8 KBs (2 KB/s) and then write 20 KBs (1 KB/s), the entire block. The writing should have taken 12 seconds but took $4 + 20 = 24$ seconds, resulting in a write speed of 0.5 KB/s, not 1 KB/s.

A direct mapping between Logical and Physical pages is not feasible inside the SSD. Therefore, each SSD has an FTL component that makes the SSD *look like an HDD*.

The **Flash Translation Layer (FTL)** is placed in the hierarchy between the **File System** and **Flash Memory**. It aims to do **three main actions**:

1. **Data Allocation and Address Translation**: It efficiently reduces Write Amplification effects (see page 29); the program pages within an erased block in order (from low to high pages), called **Log-Structured FTL**.
2. **Garbage collection**: reuse of pages with old data (Dirty/Invalid).

3. **Wear levelling:** FTL should spread across the blocks of the flash, ensuring that all of the blocks of the device wear out roughly simultaneously.

Example 3: Log-Structured FTL

Assume that a page size is 4 KB and a block consists of four pages. The write list is (`Write(pageNumber, value)`):

- `Write(100, a1)`
- `Write(101, a2)`
- `Write(2000, b1)`
- `Write(2001, b2)`
- `Write(100, c1)`
- `Write(101, c2)`

The steps are:

1. The initial state is with all pages marked as `INVALID(i)`:

Block:	0				1				2			
Page:	00	01	02	03	04	05	06	07	08	09	10	11
Content:												
State:	i	i	i	i	i	i	i	i	i	i	i	i

2. Erase block zero:

Block:	0				1				2			
Page:	00	01	02	03	04	05	06	07	08	09	10	11
Content:												
State:	E	E	E	E	i	i	i	i	i	i	i	i

3. Program pages in order and update mapping information (`Write(100, a1)`):

Table:	100 → 0												Memory
Block:	0				1				2				Flash Chip
Page:	00	01	02	03	04	05	06	07	08	09	10	11	
Content:	a1												
State:	V	E	E	E	i	i	i	i	i	i	i	i	

4. After performing four writes (`Write(100, a1)`, `Write(101, a2)`, `Write(2000, b1)`, `Write(2001, b2)`):

Table:	100 → 0 101 → 1 2000 → 2 2001 → 3												Memory
Block:	0				1				2				Flash Chip
Page:	00	01	02	03	04	05	06	07	08	09	10	11	
Content:	a1	a2	b1	b2									
State:	V	V	V	V	i	i	i	i	i	i	i	i	

5. After updating 100 and 101:

Table:	100	→	4	101	→	5	2000	→	2	2001	→	3	Memory
Block:	0				1				2				Flash Chip
Page:	00	01	02	03	04	05	06	07	08	09	10	11	
Content:	a1	a2	b1	b2	c1	c2							
State:	V	V	V	V	V	V	E	E	i	i	i	i	


When an existing page is updated, then the old data becomes obsolete. The **old versions of data are called garbage**, and (sooner or later) garbage pages must be reclaimed for new writes to take place.

The **Garbage Collection** is the **process of finding garbage blocks and reclaiming them**. It is a simple process for fully garbage blocks but more complex for partial cases.


Example 4: how garbage collection works

The steps are:

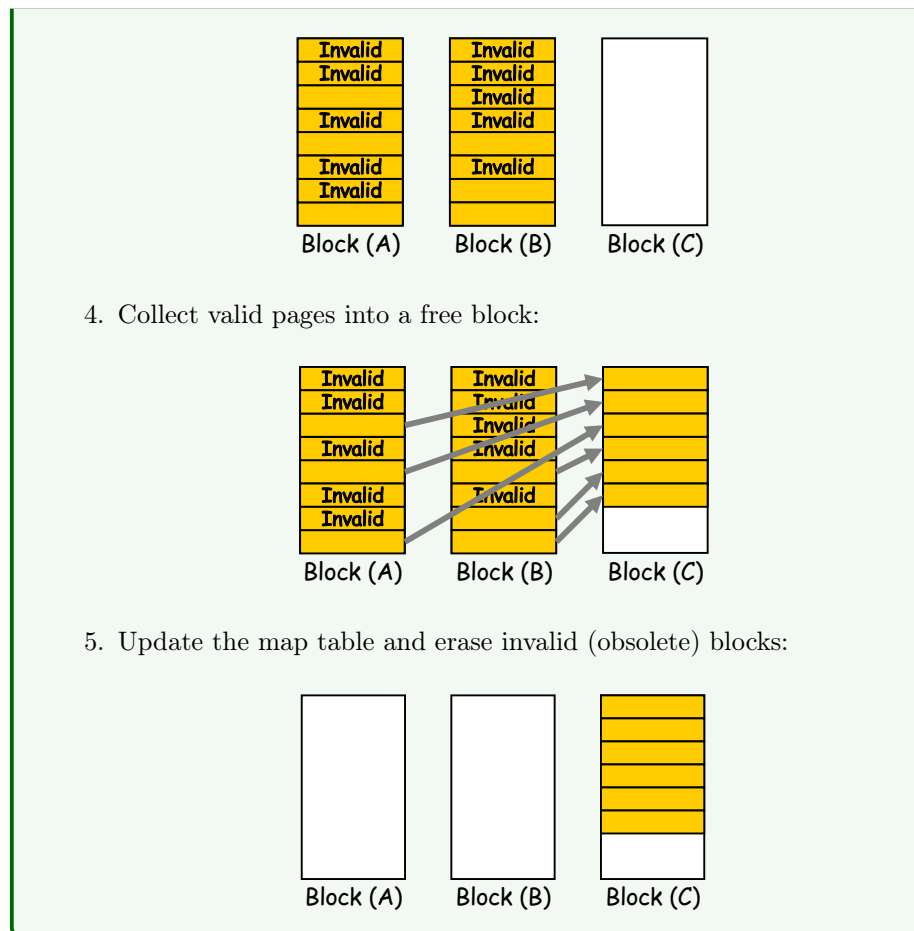
1. Update request for existing data:



2. Find a free page, and save the new data:



3. This scenario may continue until there are not enough free blocks:



⚠ Problem 1: the Garbage Collection is too expensive!

The Garbage Collection is **expensive**. It requires reading and rewriting of live data. Ideal garbage collection is a reclamation of a block that consists of only dead pages.

✓ Partial solution

Garbage Collection **costs** depend on the amount of data blocks that **must be migrated**. The **solution** to alleviate the problem is to overprovision the device by **adding extra flash capacity** (cleaning can be delayed) and **running the garbage collection in the background** using less busy periods for the disk.

⚠ Problem 2: the Ambiguity of Delete

When performing background Garbage Collection, the SSD assumes to know which pages are invalid. However, most file systems do not truly delete data. For example, on Linux, the “delete” function is `unlink()`, removing the file meta-data but not the file itself.

1. File is written on SSD
2. File is deleted
3. The Garbage Collection executes:
 - 9 pages look valid to the SSD;
 - BUT the OS knows only 2 pages are valid.

✓ Partial solution

New SSD SATA command TRIM (SCSI - UNMAP). The **OS tells the SSD that specific LBAs (page 23) are invalid and may be garbaged by the Garbage Collection.**

⚠ Problem 3: Mapping Table Size

The **size of the page-level mapping table is too large**. In fact, with a 1 TB SSD with a 4-byte entry per 4 KB page, 1 GB of DRAM is needed for mapping!

✓ Partial solution

Exists some approaches to reduce the costs of mapping:

- **Block Mapping** (block-based mapping). Mapping at block granularity to reduce the size of a mapping table. With this technique, there is a small writing problem: the FTL must read a large amount of live data from the old block and copy it into a new one.

Example 5: Block Mapping

The first four writes:

- Write(2000, a) – Write(2002, c)
- Write(2001, b) – Write(2003, d)

Table: 500 → 0

Table:	500 → 0								Memory			
Block:	0				1				2			
Page:	00	01	02	03	04	05	06	07	08	09	10	11
Content:	a	b	c	d								
State:	V	V	V	V	i	i	i	i	i	i	i	i
	Flash Chip											

Flash
Chip

And finally the last one:

- Write(2002, c')

Table:	500 → 4												Memory
Block:	0				1				2				Flash Chip
Page:	00	01	02	03	04	05	06	07	08	09	10	11	
Content:					a	b	c'	d					
State:	E	E	E	E	V	V	V	V	i	i	i	i	

- **Hybrid Mapping.** FTL maintains two tables: **log blocks** (page mapped) and **data blocks** (block mapped). The FTL will consult the page mapping table and block mapping table when looking for a particular logical block.

Example 6: Hybrid Mapping

Let's suppose the following sequence:

- Write(1000, a)
- Write(1001, b)
- Write(1002, c)
- Write(1003, d)

Log Table:													Memory
Data Table:	250 → 8												
Block:	0				1				2				Flash Chip
Page:	00	01	02	03	04	05	06	07	08	09	10	11	
Content:									a	b	c	d	
State:	i	i	i	i	i	i	i	i	V	V	V	V	

Let's update some pages:

- Write(1000, a')
- Write(1001, b')
- Write(1002, c')
- FTL updates only the page mapping information

Log Table:	1000→0	1001→1	1002→2	1003→3									Memory
Data Table:	250 →8												

Block:	0				1				2				Flash Chip
Page:	00	01	02	03	04	05	06	07	08	09	10	11	
Content:	a'	b'	c'	d'					a	b	c	d	
State:	V	V	V	V	i	i	i	i	V	V	V	V	

When needed, FTL can perform MERGE operations:

Log Table:													Memory
Data Table:	250 → 0												
Block:	0				1				2				Flash Chip
Page:	00	01	02	03	04	05	06	07	08	09	10	11	
Content:	a'	b'	c'	d'									
State:	V	V	V	V	i	i	i	i	i	i	i	i	

- **Page Mapping plus Caching.** The basic idea is to **cache the active part of the page-mapped FTL**. If a given workload only accesses a small set of pages, the translations of those pages will be stored in the FTL memory. It will perform well without high memory cost if the cache can contain the necessary working set. Cache miss overhead exists; we need to accept it.

✓ The importance of Wear Leveling

As we have mentioned, the wear leveling is essential. The erase/write cycle is limited in Flash Memory. All blocks should wear out roughly at the same time.


The log-structured approach and garbage collection help spread writing. However, a block may contain cold data: the FTL must periodically read all the live data from such blocks and re-write it elsewhere. A **disadvantage** is that the wear levelling **increases the write amplification** of the SSD and consequently decreases performance. However, to **partially fix** this, a simple policy to apply is that each flash block has an **Erase/Write Cycle Counter** and maintains the value of:

$$|\text{Max (EW cycle)} - \text{Min (EW cycle)}| < e \quad (3)$$

HDD vs SSD

Exists two metrics:

- **Unrecoverable Bit Error Ratio (UBER).** A metric for the rate of occurrence of data errors, equal to the **number of data errors per bits read**.
- **Endurance rating: Terabytes Written (TBW).** It is the total amount of data that can be written into an SSD before it is likely to fail. **The number of terabytes that may be written to the SSD while still meeting the requirements.**



2.2.2.4 RAID

RAID (Redundant Array of Independent Disks) is a data storage virtualization technology⁵ that combines multiple physical disk drive components into one or more logical units for the purposes of data redundancy, performance improvement, or both. This contrasts the previous concept of highly reliable mainframe disk drives, which were referred to as **Single Large Expensive Disks (SLED)**, also called **Just a Bunch of Disks (JBOD)** method where each disk is a separate device with a different mount point.

The data are striped across several disks accessed in parallel:

- **High data transfer rate:** large data accesses (heavy I/O operations).
- **High I/O rate:** small but frequent data accesses (light I/O operations).
- **Load Balancing** across the disks.

Two techniques exist to guarantee these features: **data striping** (improve performance) and **redundancy** (improve reliability).

Data Striping is the technique of **segmenting logically sequential data**, such as a file, so that **consecutive segments are stored on different physical storage devices**. A small quantity of terminology:

- **Striping:** data are written sequentially (a vector, a file, a table, etc) in units (stripe units such as bit, byte, and blocks) on multiple disks according to a cyclic algorithm (round robin).
- **Stripe unit:** the dimension of the data unit written on a single disk.
- **Stripe width:** number of disks considered by the striping algorithm:
 1. **Multiple independent I/O requests** will be executed in parallel by several disks, decreasing the disks' queue length (and time).
 2. Multiple disks will execute **single Multiple-Block I/O requests** in parallel, increasing the transfer rate of a single request.

The **redundancy technique is introduced because** the more physical disks in the array, the more significant the size and performance gains, but the **larger the probability of failure of a disk**.

In fact, the *probability of a failure* (assuming independent failures) in an array of 100 disks is 100 higher than the probability of a failure of a single disk! For **example**, if a disk has a **Mean Time To Failure (MTTF)** of 200.000 hours (23 years), an array of 100 disks will show an MTTF of 2000 hours (3 months).

⁵I/O virtualization: data are distributed transparently over the disks, then no action is required of the users.

The **Redundancy** is the **technique of data duplication or error correcting codes** (stored on disks different from the ones with the data) **that are computed to tolerate loss due to disk failures**. Since write operations must also update the redundant information, their **performance is worse than traditional writing**.

Data is distributed across the drives in one of several ways, referred to as **RAID levels**, depending on the required level of redundancy and performance. The different schemes, or data distribution layouts, are named by the word *RAID* followed by a number, for example RAID 0 or RAID 1. Each scheme, or RAID level, provides a different balance among the key goals: reliability, availability, performance, and capacity. The RAID levels are:

- RAID 0 striping only
- RAID 1 mirroring only
 - RAID 0+1 nested levels
 - RAID 1+0 nested levels
- RAID 2 bit interleaving (not used)
- RAID 3 byte interleaving - redundancy (parity disk)
- RAID 4 block interleaving - redundancy (parity disk)
- RAID 5 block interleaving - redundancy (parity block distributed)
- RAID 6 greater redundancy (2 failed disks are tolerated)

Note: these notes do not study the levels RAID 2 and RAID 3.

Topic	Page
RAID 0	39
RAID 1	41
RAID 0 + 1	42
RAID 1 + 0	42
RAID 4	45
RAID 5	47
RAID 6	48
Comparison and characteristics of RAID levels	49

Table 2: RAID - Table of Contents.

RAID 0

In RAID 0, the data are **written on a single logical disk and split into several blocks distributed across the disks** according to a striping algorithm.

🔍 When is it used?

It is used where **performance** and **capacity** are the primary concerns. These mean that a minimum of two drives is required.

✅ Advantages

- **Lower cost** because it does not employ redundancy (no error-correcting codes are computed and stored).
- **Best write performance** (it does not need to update redundant data and is parallelized).

⚠️ Disadvantages

Single disk failure will result in data loss.

✂️ How does it work?


The key idea is to present an **array of disks as a single large disk** and maximize parallelism by striping data across all N disks.

Now, we will give some metrics to understand how it is possible to **calculate access to a specific data block** and compare the **performance** of different RAID technologies.

To **access to a specific data blocks**, the formulas are:

$$\begin{aligned} \text{Disk} &= \text{logical_block_number} \% \text{number_of_disks} \\ \text{Offset} &= \frac{\text{logical_block_number}}{\text{number_of_disks}} \end{aligned} \quad (4)$$

For example, given the following schema:



If it is requested, the logical block is 11, and the disks are 4:

$$\begin{aligned} \text{Disk} &= 11 \% 4 = 3 \\ \text{Offset} &= 11 \div 4 = 2.75 \approx 3, \text{ then physical block 2 starting from 0} \end{aligned}$$

Note that the **chunk size is critical** because it impacts disk array performance. With **smaller chunks**, there is **greater parallelism** than with **big chunks**, which have **reduced seek times**. The typical arrays use 64 KB chunks.

To measure RAID **performance**, we focus on sequential and random workloads. Assume disks in the array have sequential transfer rate S , and the info about the disk is:

- Average seek time: 7 ms
- Average rotational delay: 3 ms
- Transfer rate: 50 MB/s

For a single large transfer (10 MB):

$$S = \frac{\text{transfer_size}}{\text{time_to_access}}$$

$$S = 10 \text{ MB} \div (7 \text{ ms} + 3 \text{ ms} + 10 \text{ MB} \div 50 \text{ MB/s}) = 47.62 \text{ MB/s}$$

If the disks in the array have a random transfer rate R , and for a set of small files (10 KB):

$$R = \frac{\text{transfer_size}}{\text{time_to_access}}$$

$$R = 10 \text{ KB} \div (7 \text{ ms} + 3 \text{ ms} + 10 \text{ KB} \div 50 \text{ MB/s}) = 0.98 \text{ MB/s}$$

Analysis

- **Capacity:** N , where N is the number of disks. Then, everything can be filled with data.
- **Reliability:** non-existent. If any drive fails, data is permanently lost. Then, the Mean Time To Failure (MTTF) equals the **Mean Time To Data Loss (MTTDL)**:

$$\text{MTTF} = \text{MTTDL}$$

- **Sequential read and write:** $N \times S$
- **Random read and write:** $N \times R$

Where N is the number of disks, S the sequential transfer rate and R the random transfer rate.

RAID 1

Although RAID 0 offers high performance, it has zero error recovery. For this reason, RAID 1 makes **two copies of all data** (again, a minimum of 2 disk drives are required).

🔍 When is it used?

It is used when **zero error recovery is not allowed**.

✓ Advantages

- **High reliability.** When a disk fails, the *second copy is used*.
- **Read of a data.** It can be *retrieved* from the *disk with shorter queueing, seek, and latency delays*.
- **Fast writes** (no error correcting code should be computed). But *still slower than standard disks* (due to duplication).

⚠ Disadvantages

- **High costs** because only 50% of the capacity is used.


✂ How does it work?

In principle, a RAID 1 can mirror the content over more than one disk. This feature gives resiliency to errors even if more than one disk breaks. Also, it allows a **voting mechanism to identify errors not reported by the disk controller**. Unfortunately, this is **never used in practice because the overhead and costs are too high**.

However, disks could be coupled if several disks are available (always in an even number). Nevertheless, the total capacity is halved, and each disk has a mirror. Then, the question is simple: *How do we organize this architecture?* The answer is the nested RAIDs: RAID 0 + 1 and RAID 1 + 0.

We define the RAID $x + y$ (or RAID xy) as:


- $n \times m$ disks in total
- m groups of n disks
- Apply RAID x to each group of n disks
- Apply RAID y considering the m groups as single disks

Figure 8: RAID $x + y$ general architecture.

The RAID 0 + 1 is a **group of striped disks (RAID 0)** that are then **mirrored (RAID 1)**. So:

1. The mirroring first (RAID 1)
2. Then the striping (RAID 0)


There are necessary almost four drives. Note that after the first failure, the model becomes a RAID 0.



The RAID 1 + 0 is a **group of mirrored disks (RAID 1)** that are then **striped (RAID 0)**. So:

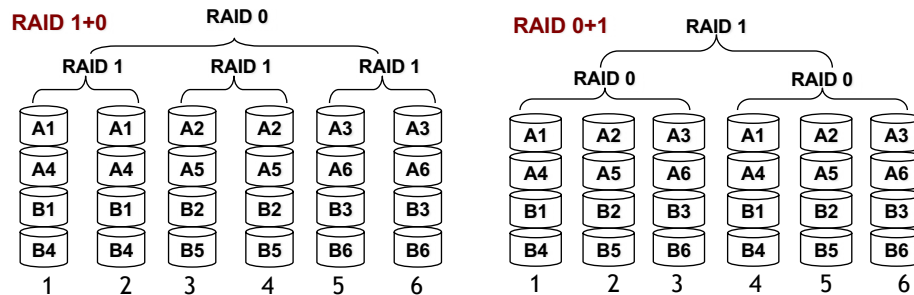
1. The striping first (RAID 0)
2. Then the mirroring (RAID 1)

There are necessary almost four drives. Usually, it is used in databases with very high workloads (fast writes).



≠ Differences between RAID 01 and RAID 10

Look at the following architectures:



What we can say is:

- The performance of RAID 01 and RAID 10 are the same.
- The **main difference is the fault tolerance**⁶!

On most implementations of RAID controllers, **RAID 01 fault tolerance is less**. Since we have only two groups of RAID 0, if two drives (one in each group) fail, the entire RAID 01 will fail. Looking at the architecture above, if Disk 1 and 4 fail, both groups will be down. So, the whole RAID 01 will fail.

On the contrary, RAID 10 since there are many groups (as the individual group is only two disks), even if three disks fail (one in each group), the RAID 10 is still functional. Looking at the architecture above, even if Disk 1, 3, and 5 fail, the RAID 10 will still be functional.

Fault Tolerance: RAID 01 \ggg RAID 10

⁶Fault tolerance is the ability of a system to maintain proper operation despite failures or faults in or more of its components.

So, given a choice between RAID 10 and RAID 01, it should be better to choose RAID 10 to have more fault tolerance.

Analysis

- **Capacity:** $N \div 2$, where N is the number of disks. Then, two copies of all data, thus half capacity.
- **Reliability:** 1 drive can fail, sometimes more! In an optimistic scenario, $N \div 2$ drives can fail without data loss.
- **Sequential write:** $(N \div 2) \times S$; two copies of all data, thus half throughput.
- **Sequential read:** $(N \div 2) \times S$; half of the read blocks are wasted, thus halving throughput.
- **Random read:** $N \times R$; it is the best scenario for RAID 1 because the read can be parallelized across all disks.
- **Random write:** $(N \div 2) \times R$; two copies of all data, thus half throughput.

Where N is the number of disks, S is the sequential transfer rate, and R is the random transfer rate.

It is essential to observe RAID 1. There is a notorious **problem** called the **Consistent Update Problem**.

Mirrored writes should be atomic. Then, all copies are written, or none are written. Unfortunately, this is very **difficult to guarantee** sometimes, for example, in a power failure scenario. To **solve** this problem, many RAID controllers include a **write-ahead log**, a **battery backend**, and **non-volatile storage of pending writes**. With this system, a **recovery procedure ensures the recovery of the out-of-sync mirrored copies**.

RAID 4

RAID 4 consists of **block-level striping with a dedicated parity disk**.

? When is it used?

RAID 1 offers highly reliable data storage, but it uses half the space of the array capacity. To achieve the **same level of reliability without wasting capacity**, it is possible to use RAID 4, which uses **information coding techniques to build lightweight error recovery mechanisms**.

✓ Advantages

- **Good performance** of random reads.

🚫 Disadvantages

- **Random Write performance is terrible** due to being *bottlenecked* by the parity drive.

✂ How does it work?

Disk N only stores parity information for the other $N - 1$ disks. The parity is calculated using the XOR operation⁷.

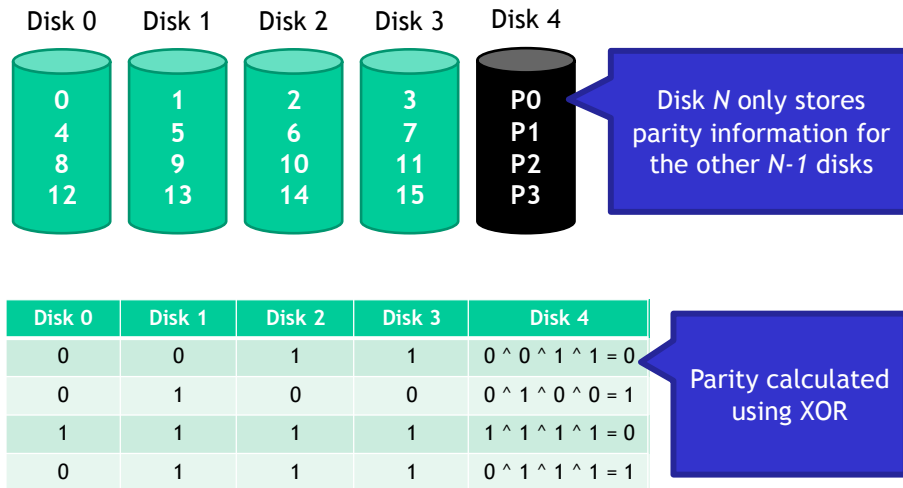


Figure 9: RAID 4 - *How does it work?*

The **serial** or **random read** is not a problem in RAID 4 because there is a **parallelization across all non-parity blocks** in the stripe despite a tiny performance reduction due to the parity disk.

⁷“A or B, but not A and B”

During the parity writing, the blocks are updated. The **random write** performance is affected by the approach used:

- **Additive parity** (as known as reconstruct-writes):
 1. Read all other blocks in a stripe in parallel;
 2. XOR those with the new block to form a new parity block;
 3. Write the new data block and new parity block to disks.

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	0	0	1	$0 \wedge 0 \wedge 0 \wedge 1 = 1$

- **Subtractive parity** (as known as read-modify-writes):
 1. Read only the old data block to be updated and the old parity block;
 2. Compute the new parity block using:

$$P_{new} = (D_{new} \vee D_{old}) \vee P_{old}$$

Where \vee is the logical XOR.

3. Write the new data block and new parity block to disks.

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	0	1 0	1	$0 \wedge 0 \wedge 1 \wedge 1 = 1$

Despite the **sequential write** does **not suffer any performance effect** from the parity disk. Because it uses full-stripe write:

1. Buffer all data blocks of a stripe
2. Compute the parity block
3. Write all data and parity blocks in parallel [5]

Analysis

- **Capacity:** $N - 1$, where N is the number of disks, and the -1 is because one is dedicated to the parity disk.
- **Reliability:** 1 drive can fail. Massive performance degradation during partial outage.
- **Sequential read/write:** $(N - 1) \times S$; parallelization across all non-parity blocks in the stripe (parity disk has no effect).
- **Random read:** $(N - 1) \times R$; reads parallelize over all but the parity drive (parity disk has no effect).
- **Random write:** $R \div 2$; writes serialize due to the parity drive, and each write requires one read and one write of the parity drive.

Where N is the number of disks, S is the sequential transfer rate, and R is the random transfer rate.

RAID 5

RAID 5 rotates parity blocks across stripes.

? When is it used?

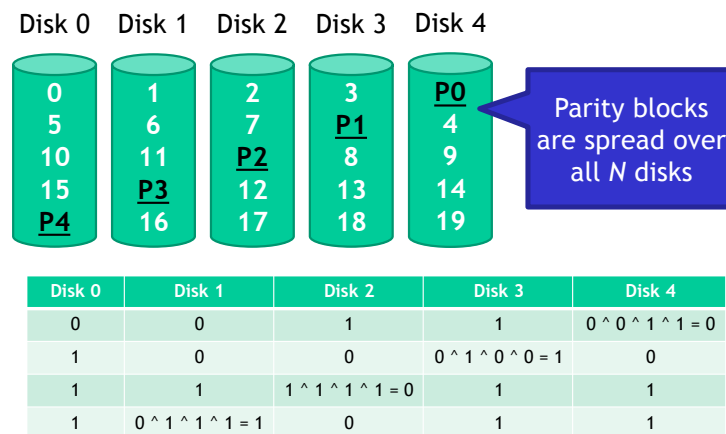
Unlike in RAID 4, parity information is distributed among the drives. This technique is **used to improve significantly the random write throughput** against the RAID 4 system.

✓ Advantages

- Improved random write despite the RAID 4 system.

✂ How does it work?

The writes are spread roughly evenly across all drives.



The **random write** in RAID 5 is:

1. Read the target block and the parity block
2. Use subtraction to calculate the new parity block
3. Write the target block and the parity block

Thus, **four total operations** (two reads, two writes) **are distributed across all drives**.

📊 Analysis

- **Capacity:** $N - 1$ (same as RAID 4), where N is the number of disks.
- **Reliability:** 1 drive can fail (same as RAID 4). Massive performance degradation during partial outage.

- **Sequential read/write:** $(N - 1) * S$ (**same as RAID 4**); parallelization across all non-parity blocks in the stripe (parity disk has no effect).
- **Random read:** $N \times R$; unlike RAID 4, **reads parallelize over all drives**.
- **Random write:** $(N \div 4) \times R$; unlike RAID 4, **writes parallelize over all drives**, and each write requires two reads and two writes, hence $N \div 4$.

Where N is the number of disks, S is the sequential transfer rate, and R is the random transfer rate.


RAID 6

RAID 6 can tolerate multiple disk faults (**more fault tolerance**) concerning RAID 5. It tolerates two concurrent failures.

✂ How does it work?

It uses Solomon-Reeds codes with two redundancy schemes and requires $N + 2$ disks (where N is the number of disks). Note that the **minimum set is 4 data disks**.

Unfortunately, it has a **high overhead for writes** (computation of parities) because each write requires six disk accesses due to the need to update both the P and Q parity blocks (slow writes).



Comparison and characteristics of RAID levels

The following table shows eight fundamental properties of the RAID levels.

- N : number of drives
- R : random access speed
- S : sequential access speed
- D : latency to access a single disk

	RAID 0	RAID 1	RAID 4	RAID 5
Capacity	N	$N \div 2$	$N - 1$	$N - 1$
Reliability	0	$1 \vee : N \div 2$	1	1
Sequential Read	$N \times S$	$(N \div 2) \times S$	$(N - 1) \times S$	$(N - 1) \times S$
Sequential Write	$N \times S$	$(N \div 2) \times S$	$(N - 1) \times S$	$(N - 1) \times S$
Random Read	$N \times R$	$N \times R$	$(N - 1) \times R$	$N \times R$
Random Write	$N \times R$	$(N \div 2) \times R$	$R \div 2$	$(N \div 4) \times R$
Read	D	D	D	D
Write	D	D	$2 \times D$	$2 \times D$

Table 3: Comparison of RAID levels.

Where the throughput is:

- Sequential Read
- Sequential Write
- Random Read
- Random Write

And the latency is:

- Read
- Write

RAID level	Capacity	Reliability	R/W performance	Rebuild performance	Suggested applications
0	100%	N/A	Very good	Good	Non critical data
1	50%	Excellent	Very good / good	good	Critical information
5	$(n-1)/n$	Good	Good/ fair	Poor	Database, transaction based applications
6	$(n-2)/n$	Excellent	Very good/ poor	Poor	Critical information, w/minimal
1+0	50%	Excellent	Very good/ good	Good	Critical information, w/better performance

Figure 10: Characteristics of RAID levels.

RAID 0 has the best performance and most capacity.

RAID 1 (10 better than 01) or **RAID 6** have the greatest error recovery.

RAID 5 has the better *balance* between space, performance, and recoverability.

2.2.3 Networking (architecture and technology)

3 Software Infrastructure

3.1 Virtualization

3.2 Computing Architectures

4 Methods

4.1 Reliability and availability of datacenters

4.2 Disk performance

4.2.1 HDD

We can calculate some performance metrics related to the types of delay of HDD (page 27).

- **Full Rotation Delay** R is:

$$R = \frac{1}{\text{DiskRPM}} \quad (5)$$

And in seconds:

$$R_{\text{sec}} = 60 \times R \quad (6)$$

From the R_{sec} we can also calculate the **total rotation average**:

$$T_{\text{rotation AVG}} = \frac{R_{\text{sec}}}{2} \quad (7)$$

- **Seek Time**, the **time to move the head to a different track**, which is divided into several phases:
 - Acceleration
 - Coasting (constant speed)
 - Deceleration
 - Settling

The T_{seek} modelling considers a linear dependency with the distance. Also, the **seek average** is:

$$T_{\text{seek AVG}} = \frac{T_{\text{seek MAX}}}{3} \quad (8)$$

- **Transfer time**. It is the **time that data is either read from or written to the surface**. It includes the time the **head needs to pass on the sectors** and the **I/O transfer**:

$$T_{\text{transfer}} = \frac{\text{R/W of a sector}}{\text{Data transfer rate}} \quad (9)$$

The **Controller Overhead** is the **buffer management** (data transfer) and **interrupt sending time**.

Transfer time and Controller Overhead are together because they are required to calculate some interesting metrics.

- **Service Time** $T_{\text{I/O}}$

$$T_{\text{I/O}} = T_{\text{seek}} + T_{\text{rotation}} + T_{\text{transfer}} + T_{\text{overhead}} \quad (10)$$

- **Response Time**

$$T_{\text{queue}} + T_{\text{I/O}} \quad (11)$$

Where T_{queue} depends on queue-length, resource utilization, mean and variance of disk service time and request arrival distribution.




Figure 11: Service and response time.

Exercise 1: mean service time of an I/O operation

The data of the exercise are:

- Read/Write of a sector of 512 bytes (0.5 KB)
- Data transfer rate: 50 MB/sec
- Rotation speed: 10000 RPM (Round Per Minute)
- Mean seek time: 6 ms
- Overhead Controller: 0.2 ms

To calculate the *service time* $T_{I/O}$, we need the following information:

- T_{seek} , which we already have, and it is 6 ms.
- $T_{rotation}$
- $T_{transfer}$
- $T_{overhead}$, which we already have, and it is 0.2 ms.

We also know the rotation and transfer information, but we want to know the *mean* service time. Then we calculate the total rotation average $T_{rotation\ AVG}$:

$$\begin{aligned}
 R &= \frac{1}{DiskRPM} = \frac{1}{10000} = 0.0001 \\
 R_{sec} &= 60 \cdot R = 60 \cdot 0.0001 = 0.006 \text{ seconds} \\
 T_{rotation\ AVG} &= \frac{R_{sec}}{2} = \frac{0.006}{2} = 0.003 \text{ seconds} = 3 \text{ ms}
 \end{aligned}$$

Finally, the transfer time is easy to calculate because we have the R/W of a sector and the data transfer rate. First we do a conversions from

megabytes to kilobytes:

$$\begin{aligned}
 \text{Data transfer rate:} &= 50 \text{ MB/sec} \\
 &= 50 \cdot 1024 \text{ KB/sec} \\
 &= 51200 \text{ KB/sec} \\
 T_{\text{transfer}} &= \frac{0.5 \text{ KB/sec}}{51200 \text{ KB/sec}} \\
 &= 0.000009765625 \text{ sec} \cdot 1000 \\
 &= 0.009765625 \text{ ms} \approx 0.01 \text{ ms}
 \end{aligned}$$

The exercise can be completed by calculating the mean I/O service time required:

$$\begin{aligned}
 T_{\text{I/O}} &= T_{\text{seek}} + T_{\text{rotation}} + T_{\text{transfer}} + T_{\text{overhead}} \\
 T_{\text{I/O}} &= 6 + 3 + 0.01 + 0.2 = 9.21 \text{ ms}
 \end{aligned}$$

The previous service time is supposed to be only for very **pessimistic cases** where **sectors are fragmented on the disk in the worst possible way**. This can happen because the files are tiny (each file is contained in one block) or the disk is externally fragmented.

Thus, each access to a sector requires to pay rotational latency and seek time. This is not the case in many circumstances because the files are larger than one block and stored contiguously.

We can measure the **Data Locality DL** of a disk as the **percentage of blocks that do not need seek or rotational latency to be found**.

Thanks to the Data Locality, it is possible to calculate the **Average Service Time**:

$$T_{\text{I/O AVG}} = (1 - DL) \cdot (T_{\text{seek}} + T_{\text{rotation}}) + T_{\text{transfer}} + T_{\text{controller}} \quad (12)$$

Exercise 2: data locality

The data of the exercise are:

- Read/Write of a sector of 512 bytes (0.5 KB)
- Data Locality: $DL = 75\%$
- Data transfer rate: 50 MB/sec
- Rotation speed: 10000 RPM (Round Per Minute)
- Mean seek time: 6 ms
- Overhead Controller: 0.2 ms

Since the Data Locality is 75%, only 25% of the operations are affected by the DL:

$$(1 - DL) = (1 - 0.75) = 0.25$$

See the exercise on page 56 to understand the values of T_{seek} , T_{rotation} , T_{transfer} and T_{overhead} :

- $T_{\text{seek}} = 6$
- $T_{\text{rotation}} = 3$
- $T_{\text{transfer}} = 0.01$
- $T_{\text{overhead}} = 0.2$

Finally the average time for read/write a sector of 0.5 KB with a DL of 75% is:

$$\begin{aligned} T_{\text{I/O AVG}} &= 0.25 \cdot (6 + 3) + 0.01 + 0.2 \\ &= 0.25 \cdot 9 + 0.21 \\ &= 2.46 \text{ ms} \end{aligned}$$

Exercise 3: influence of “not optimal” data allocation

The data of the exercise are:

- 10 blocks of 1/10 MB for each block (10 blocks of 1/10 MB “not well” distributed on disk)
- $T_{\text{seek}} = 6 \text{ ms}$
- $T_{\text{rotation AVG}} = 3 \text{ ms}$
- Data transfer rate: 50 MB/sec

In the exercise you were asked to calculate the time taken to transfer a 1 MB file with 100% and 0% data locality:

- Data Locality equals to 100%:
 - An initial seek (6 ms)
 - A total rotation average (3 ms)
 - Now it’s possible to do the 1MB global transfer directly because there are no blocks to seek or rotation latency:

$$1 \text{ MB of } 50 \text{ MB} = \frac{1}{50} = 0.02 \text{ seconds} \cdot 1000 = 20 \text{ ms}$$

- The total time is:

$$T = 6 + 3 + 20 = 29 \text{ ms}$$

- Data Locality equals to 0%:
 - An initial seek (6 ms)
 - A total rotation average (3 ms)
 - In this case, it's not possible to do a global transfer directly, because each block is affected by the seek or rotation latency. Then we have to transfer block by block and calculate the delay:

$$1 \text{ MB of } 10 \text{ MB} = \frac{1}{10} = 0.1 \text{ seconds} \cdot 1000 = 100 \text{ ms}$$

- The total time is:

$$T = (6 + 3 + 2) \cdot 10 = 110 \text{ ms}$$

Where 10 is the number of blocks.

Note: the controller times is not considered.

4.3 Scalability and performance of datanceters

References

- [1] L.A. Barroso, U. Hölzle, and P. Ranganathan. *The Datacenter as a Computer: Designing Warehouse-Scale Machines, Third Edition*. Synthesis Lectures on Computer Architecture. Springer International Publishing, 2022.
- [2] E.D. Lazowska. *Quantitative System Performance: Computer System Analysis Using Queueing Network Models*. Prentice-Hall, 1984.
- [3] Jacopo Marino and Fulvio Rizzo. Is the computing continuum already here?, 2023.
- [4] Gianluca Palermo. Lesson 1, computing infrastructures. Slides from the HPC-E master’s degree course on Politecnico di Milano, 2024.
- [5] Ming-Chang Yang. Lesson 2, raid and data integrity. Slides from the CSCI5550 Advanced File and Storage Systems course on the Chinese University of Hong Kong, 2020.

Index

A

actual size on the disk a	24
Additive parity	46
Availability Zones (AZs)	12
Average Service Time	57

B

Blade Servers	18
Block Mapping	34

C

C-LOOK	28
C-SCAN (Circular SCAN)	28
cache	27
Clusters	23
Computing Continuum	5
Computing Infrastructure	4
Computing Regions (CRs)	12
Consistent Update Problem	44
Controller Overhead	27, 55

D

data blocks	23
Data Center	4
Data Locality DL	57
Data Striping	37
Dirty Page	29

E

Edge Computing	8
Embedded System	7
Empty Page	29
Endurance rating: Terabytes Written (TBW)	36
Erase Block	29
Erase/Write Cycle Counter	36
Erased Page	29
external fragmentation	25

F

First Come, First Serve (FCFC)	28
Flash Translation Layer (FTL)	30
Fog Computing	8
Full Rotation Delay	55

G

Garbage Collection	32
Geographic Areas (GAs)	12

H

Hard Disk Drive (HDD)	22, 26
-----------------------	--------

Hybrid Mapping	35
I	
In Use Page	29
internal fragmentation	24
Internet of Things (IoT)	7
Invalid Page	29
J	
Just a Bunch of Disks (JBOD)	37
L	
LBA (Logical Block Address)	23
Log-Structured FTL	30
M	
Mean Time To Data Loss (MTTDL)	40
Mean Time To Failure (MTTF)	37
N	
non-volatile memory express (NVMe)	22
P	
Page	29
Page Mapping plus Caching	36
PCIe (peripheral component interconnect express)	22
R	
Rack Servers	17
Rack Units	17
RAID (Redundant Array of Independent Disks)	37
RAID levels	38
Read caching	27
read-modify-writes	46
reconstruct-writes	46
Redundancy	38
Response Time	55
Revolutions Per Minute (RPM)	26
Rotational Delay	27
S	
SCAN	28
seek average	55
Seek Delay	27
Seek Time	55
Server	15
Service Time	55
Shortest Seek Time First (SSTF)	28
Single Large Expensive Disks (SLED)	37
solid-state drive (SSD)	22, 29
Stripe unit	37

Stripe width	37
Striping	37
Subtractive parity	46
T	
torn write	26
total rotation average	55
Tower Server	16
track buffer	27
Transfer time	27, 55
U	
Unrecoverable Bit Error Ratio (UBER)	36
V	
Valid Page	29
W	
Warehouse-Scale Computers (WSCs)	11
wasted disk space w	24
Write amplification (WA)	29
Write Back cache	27
Write caching	27
Write Through cache	27
write-ahead log	44